

Exploratory Data Analysis of Airbnb Listing

Art Tay

Introduction

Airbnb, Inc. is a software company that operates a market place for people to rent out their properties on an ad hoc basis ([Wikipedia](#)). In specific dataset addressed in this report comes from Inside Airbnb, “a mission driven project that provides data ... about Airbnb’s impact of residential communities.” ([Website](#)). The data set contains 12 predictive variables: 4 categorical and 8 numeric. All of the roughly 40,000 listings in this particular dataset were located in New York City. The ultimate goal of this analysis was to categorize the different types of Airbnb listings and their hosts, with a particular emphasis on determine which listings had the most business traffic.

Summary of Results

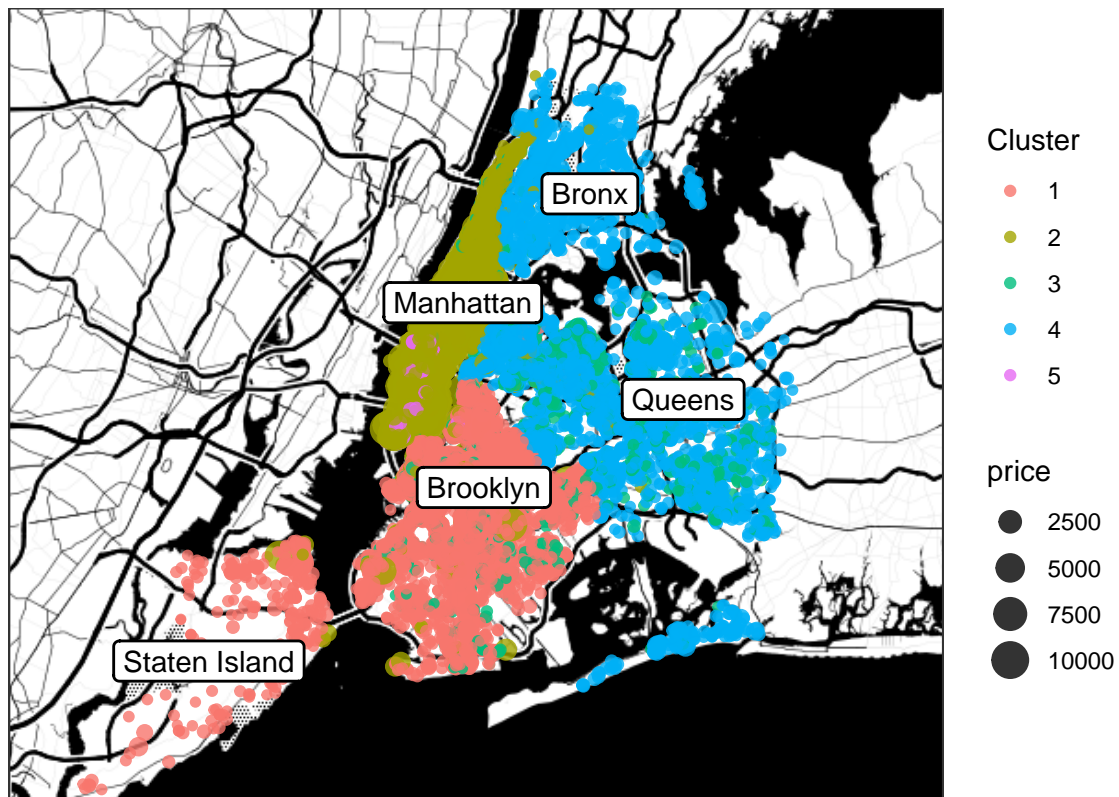


Figure 1: Geographic plot of cluster and price.

Ultimately 5 different groups of Airbnbs were identified. Groups 1, 2, and 4 were mainly geographically tied, representing Brooklyn, Manhattan, and Queens respectively. Group 5 is made up of the most expensive listings. They are all located in Manhattan, but are rarely actually booked. Finally, Group 3 was determined to be the busiest Airbnbs. Group 3 was geographically dispersed, medianly priced, and had the most number and recency of reviews. If we assume that every review represents at least one booking, then Group 3 is by far the listings that attract the most business.

Methodology

Data Cleaning

Before the modeling process could begin 3 important issues had to be addressed in the raw data: missing values, feature engineering, and skewness.

Missing Values

Upon first inspection there are only NA values in the `reviews_per_month` variable; however, there are also blanks in `last_review`, and zero values that appear erroneous. Any zeros in `price`, `latitude`, `longitude`, and `minimum_night` was replaced with an NA. Furthermore, any blank spaces in a factor variable was also considered to be NA. Below is a plot of the pattern and quantity of missing values after recoding.

Although it appears that there are roughly 10,000 missing values for `review_per_month` and `last_review` the pattern is not random, in fact, **Figure 1** indicates that these two variables are always missing together. Upon further inspection, the missingness of these two variables perfectly correlates with `number_of_review` being equal to 0. Ultimately these are not actually missing values, they simply indicate that a particular listing has never been reviewed.

There were 11 listings where the `price` was 0. Because price is such a key factor in the desirability of a particular listing, these 11 values were imputed using the k-nearest neighbors algorithm, where $k = 10$, given the large number of observations. Since `price` is numeric the median of the neighbor was chosen.

Feature Engineering

After the nature of missing data was determined, related data cleaning commenced. First, NAs for `review_per_month` were replaced with 0. Next, `last_review` was replaced by `last_review_year` to reduce the number of factor levels. In hindsight, the day and month could have also been extracted as new features. Then, the `neighbourhood` factor was dropped because of the copious number of levels and the majority of the information being captured in `neighbourhood_group`. Perhaps a hierarchical feature could have been engineered. Finally, all factor variables were converted into dummy variables and all numeric variables were normalized. This step was a requirement of the imputation algorithm used. Note that `impStatus` was added as a binary factor 1 if price was imputed, and 0 otherwise.

Skewness

Skewness was a key issue among also every single numeric variable.

The `price` was transformed using a log function, since all zero values were treated as NA. This was done to help the imputation process. The decision was made to not transform any of the other numeric predictors before clustering as completed. There are competing views in the literature about whether or not clustering algorithms are affected by skewness either negatively or positively. Because the clusters were ultimately fit using euclidean distances and inter-cluster variability, applying transformations that would artificially reduce variance seem incorrect. Post clustering a Yeo-Johnson transformation was applied to the numeric variables




Figure 2: Quantity and Pattern of Missing Data. The blue bars represent the quantity of the pattern shown of the matrix below. The black bars show the number of missing values of the variable to the right across all patterns.

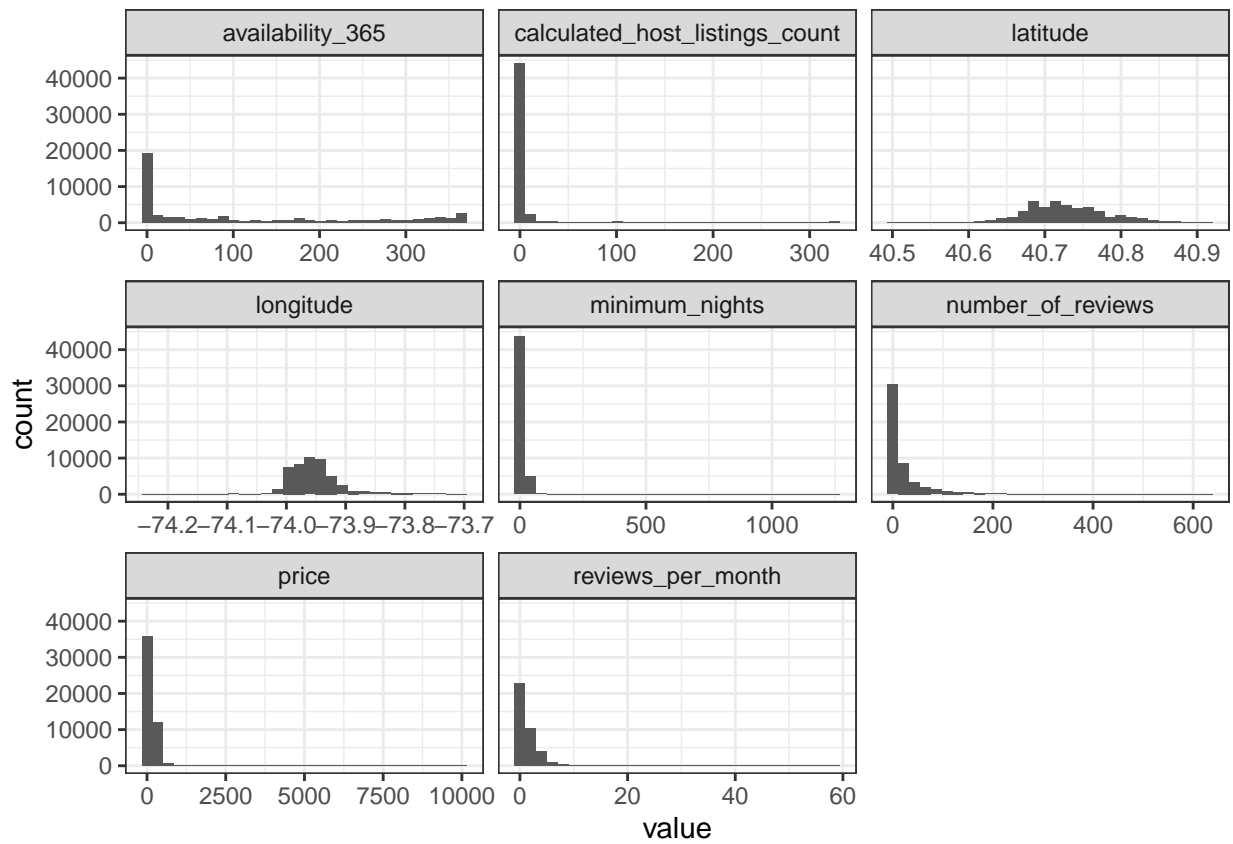


Figure 3: Histogram of raw numeric features before any transformations. Most distribution are clearly right skewed

only to increase visibility of plot features. Yeo-Johnson was chosen over Box-Cox because it can work with normalized data.

Cluster Modeling

Cluster was applied to the entire cleaned and imputed dataset as outlined above. The cluster algorithm used was implemented by the `NBClust` package. Clusters were determined based on a minimization of total within-cluster variance where dissimilarity was defined as Euclidean distance. This is known as WARD clustering (Ward 1963). The number of clusters was determined based on maximizing the “silhouette” index (Rousseeuw, 1987). 2-15 clusters were tried.

Result

The cluster modeling determined that 5 was the optimal number of clusters.

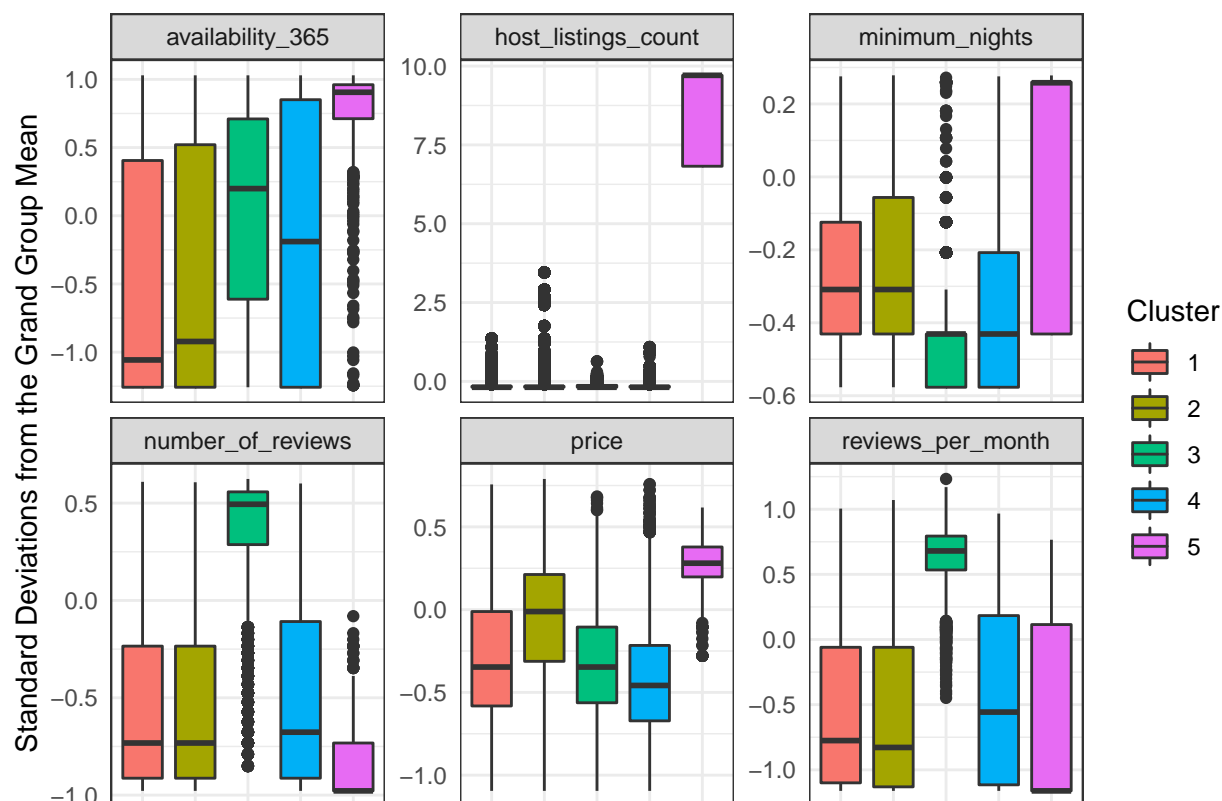


Figure 4: Histogram of raw numeric features before any transformations. Most distribution are clearly right skewed

Appendix-Code




Figure 5: Histogram of raw numeric features before any transformations. Most distribution are clearly right skewed




Figure 6: Histogram of raw numeric features before any transformations. Most distribution are clearly right skewed

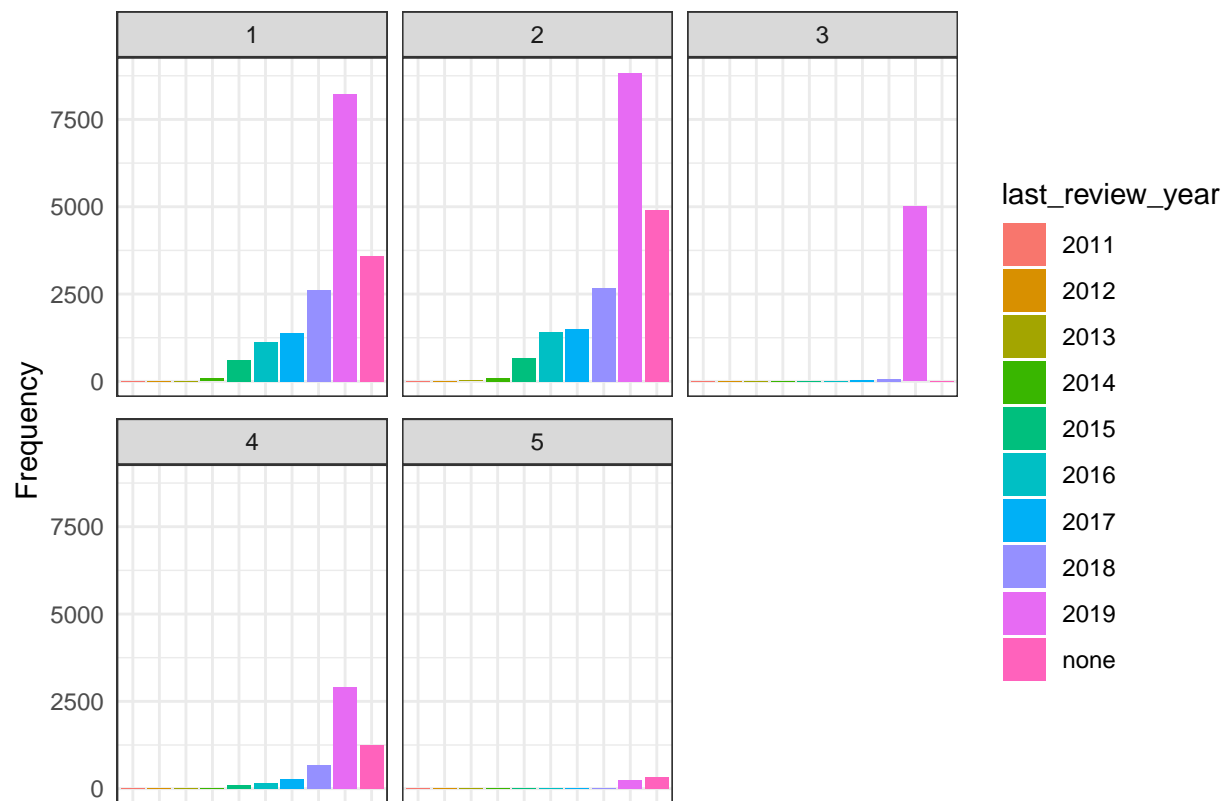


Figure 7: Histogram of raw numeric features before any transformations. Most distribution are clearly right skewed


```

# Libraries
library(VIM)
library(naniar)
library(tidyverse)
library(fastDummies)
library(NbClust)
library(recipes)
library(corrplot)

# Load in Data
data_full <- read.csv("AB_NYC_2019.csv", stringsAsFactors = T, header = T)

```

Data Cleaning

```

# Removing uninformative variables (names).
data_quant <- data_full %>% select(-c(id, host_id, name, host_name))

```

Missing Data

```

# Code value that might mean missing.
# price == 0 -> NA
# latitude == 0 -> NA
# longitude == 0 -> NA
# min_night == 0 -> NA

data_quant_mis <- data_quant %>%
  mutate(price = ifelse(price == 0, NA, price)) %>%
  mutate(latitude = ifelse(latitude == 0, NA, latitude)) %>%
  mutate(longitude = ifelse(longitude == 0, NA, longitude)) %>%
  mutate(minimum_nights = ifelse(minimum_nights == 0, NA, minimum_nights)) %>%
  # It was done using replace to maintain factor coding
  mutate(room_type = replace(room_type,
    room_type == "" | room_type == " ", NA)) %>%
  mutate(neighbourhood_group = replace(neighbourhood_group,
    neighbourhood_group == "" | neighbourhood_group == " ", NA)) %>%
  mutate(neighbourhood = replace(neighbourhood,
    neighbourhood == "" | neighbourhood == " ", NA)) %>%
  mutate(last_review = replace(last_review,
    last_review == "" | room_type == " ", NA))

save(data_quant_mis, file = "Report Figures/data_quant_mis.rds")

# Plot the pattern of missing values
miss_pattern_plot <- gg_miss_upset(data_quant_mis, nintersects = NA, text.scale = 2,
  mainbar.y.label = "Count of Observations with Missing Data",
  point.size = 5, line.size = 2, matrix.color = "orange",
  main.bar.color = "skyblue")

```

Feature Engineering

```
# Handles all missing values that are not price
data_quant_fe <- data_quant_mis %>%
  # Change NA reviews per month to be 0.
  mutate(reviews_per_month =
    ifelse(is.na(reviews_per_month), 0, reviews_per_month))%>%
  # Create a new variable last_review_year to reduce dimensionality.
  mutate(last_review_year = substring(last_review, 1, 4)) %>%
  # Modify NA last_review to be a new level "none".
  mutate(last_review_year = replace(last_review_year,
    is.na(last_review_year), "none")) %>%
  # Cast to be factor.
  mutate(last_review_year = as.factor(last_review_year)) %>%
  # Remove old last review variable
  select(-last_review) %>%
  # natural log transform price.
  mutate(price = log(price))
```

```
# Centering a scale numerics.
data_quant_mis_numeric <- data_quant_fe %>% select(where(is.numeric))

# Retain vector to enable back transformation.
means <- apply(data_quant_mis_numeric, MARGIN = 2, FUN = mean, na.rm = T)
sds <- apply(data_quant_mis_numeric, MARGIN = 2, FUN = sd, na.rm = T)

# Apply transformation
for(i in 1:length(means)){
  data_quant_mis_numeric[, i] <-
    (data_quant_mis_numeric[, i] - means[i]) / sds[i]
}
```

```
# Dummy variables (one-hot-encoding)
data_quant_mis_factor <- data_quant_fe %>% select(where(is.factor)) %>%
  # too many levels.
  select(-neighbourhood)

data_quant_mis_factor <- dummy_cols(data_quant_mis_factor,
  remove_selected_columns = T)

# Combine cleaned data set for imputation.
data_quant_clean <- cbind(data_quant_mis_numeric, data_quant_mis_factor)
```

```
# kNN imputation with k = 10
data_quant_clean_imputed <- VIM::kNN(as.matrix(data_quant_clean), k = 10)

# Create a factor based on whether or not price as imputed
impStatus <- as.numeric(data_quant_clean_imputed$price_imp)
```

```
# get the original columns
data_quant_clean_imputed <- data_quant_clean_imputed[, 1:26]
data_quant_clean_imputed$impStatus <- impStatus
```

Imputation

Skewness

```
# Histogram of all numerics
plot_skew <- data_quant %>%
  select(where(is.numeric)) %>%
  pivot_longer(cols = everything()) %>%
  ggplot(aes(value)) +
  geom_histogram() +
  facet_wrap(~name, scales = "free_x") +
  theme_bw()

save(plot_skew, file = "Report Figures/plot_skew.rds")
```

Clustering

```
#data_clusters <- NbClust(data = data_quant_clean_imputed,
  #distance = "euclidean",
  #method = "ward.D",
  #index = "silhouette")
```

```
#save(data_clusters, file = "data_clusters.rds")
load(file = "data_clusters.rds")
```

Cluster 1 cleaning

```
# Back Transform numeric data.
# Extract numeric columns.
data_imputed_numeric <- data_quant_clean_imputed %>%
  select(1:8)

# Undo centering and scaling.
for(i in 1:length(means)){
  data_imputed_numeric[, i] <-
    data_imputed_numeric[, i] * sds[i] + means[i]
}

# Undo log transformation on price.
data_imputed_numeric <- data_imputed_numeric %>%
  mutate(price = exp(price)) %>%
  # fixes rounding error
  mutate(availability_365 = round(availability_365))
```

```

# Extract original factor formats.
data_imputed_factor <- data_quant_fe %>% select(where(is.factor))

# Confirm no missing
#sum(is.na(data_imputed_factor))

data_clean_cluster_1 <- cbind(data_imputed_numeric,
                              data_imputed_factor,
                              "impStatus" = as.factor(
                                data_quant_clean_imputed$impStatus),
                              "Cluster" = as.factor(
                                data_clusters$Best.partition))

```

Cluster Statistics and Plots

```

cluster_summary_numeric <- data_clean_cluster_1 %>% group_by(Cluster) %>%
  summarise(across(where(is.numeric),
                    list(median = median, sd = sd)))

cluster_summary_factor <- data_clean_cluster_1 %>%
  select(where(is.factor)) %>%
  select(-impStatus) %>%
  select(Cluster, neighbourhood_group) %>%
  pivot_longer(!Cluster) %>%
  count(Cluster, name, value, .drop = F, sort = T)

```

```

plot_neighborhoods <- cluster_summary_factor %>%
  ggplot(aes(x = value, y = n, fill = value)) +
  geom_bar(stat = 'identity') +
  facet_wrap(~Cluster, scales = "free_x") +
  labs(x = "", y = "Frequency",
       fill = "neighbourhood_group") +
  theme_bw() +
  theme(axis.text.x = element_blank(),
        axis.ticks = element_blank())

save(plot_neighborhoods, file = "Report Figures/plot_neighborhoods.rds")

```

```

# Plot of room type frequency by cluster.
cluster_summary_type <- data_clean_cluster_1 %>%
  select(where(is.factor)) %>%
  select(-impStatus) %>%
  select(Cluster, room_type) %>%
  pivot_longer(!Cluster) %>%
  count(Cluster, name, value, .drop = F, sort = T)

```

```

plot_type <- cluster_summary_type %>%
  ggplot(aes(x = value, y = n, fill = value)) +
  geom_bar(stat = 'identity') +
  facet_wrap(~Cluster, scales = "free_x") +
  labs(x = "", y = "Frequency",
       fill = "room_type") +
  theme_bw() +
  theme(axis.text.x = element_blank(),
        axis.ticks = element_blank())

save(plot_type, file = "Report Figures/plot_type.rds")

```

```

# Plot last_review_year frequency by cluster.
cluster_summary_year <- data_clean_cluster_1 %>%
  select(where(is.factor)) %>%
  select(-impStatus) %>%
  select(Cluster, last_review_year) %>%
  pivot_longer(!Cluster) %>%
  count(Cluster, name, value, .drop = F, sort = T)

plot_year <- cluster_summary_year %>%
  ggplot(aes(x = value, y = n, fill = value)) +
  geom_bar(stat = 'identity') +
  facet_wrap(~Cluster, scales = "free_x") +
  labs(x = "", y = "Frequency",
       fill = "last_review_year") +
  theme_bw() +
  theme(axis.text.x = element_blank(),
        axis.ticks = element_blank())

save(plot_year, file = "Report Figures/plot_year.rds")

```

Factor Plot

```

remove_skew <- recipe(~., x = data_imputed_numeric) %>%
  step_normalize(all_predictors()) %>%
  step_YeoJohnson(all_predictors()) %>%
  prep(retain = T)

data_noskew_cluster_1 <- bake(remove_skew, new_data = NULL)

data_noskew_cluster_1 <- cbind(data_noskew_cluster_1,
                               data_imputed_factor,
                               "impStatus" = as.factor(
                                 data_quant_clean_imputed$impStatus),
                               "Cluster" = as.factor(
                                 data_clusters$Best.partition))

```

```

numeric_plot_noskew_1 <- data_noskew_cluster_1 %>%
  select(price, number_of_reviews,
         reviews_per_month,
         calculated_host_listings_count,
         availability_365, Cluster, minimum_nights) %>%
  rename(host_listings_count = calculated_host_listings_count) %>%
  pivot_longer(!Cluster) %>%
  ggplot(aes(y = value, x = Cluster, fill = Cluster)) +
  geom_boxplot() +
  facet_wrap(~name, scale = "free_y") +
  labs(x = "", y = "Standard Deviations from the Grand Group Mean") +
  theme_bw() +
  theme(axis.text.x = element_blank(),
        axis.ticks = element_blank())

save(numeric_plot_noskew_1, file = "Report Figures/numeric_plot_noskew_1.rds")

```

Numeric Plots

```

NYC_Map <- get_stamenmap(
  bbox = c(left = -74.3, bottom = 40.5, right = -73.6, top = 40.95),
  maptype = "toner-background", color = "bw")

```

```

save(NYC_Map, file = "Report Figures/NYC_Map.rds")
#load("Report Figures/NYC_Map.rds")

```

```

df_neighbourhood_labels <- data.frame(
  longitude = c(-73.8648, -73.9442, -73.96, -73.7949, -74.1502),
  latitude = c(40.8448, 40.6782, 40.7831, 40.7282, 40.5795),
  borough = c("Bronx", "Brooklyn", "Manhattan", "Queens", "Staten Island")
)

```

```

map_plot <- ggmap(NYC_Map) + geom_point(data = data_clean_cluster_1,
  mapping = aes(x = longitude, y = latitude,
                col = Cluster, size = price),
  alpha = .8) +
  geom_label(data = df_neighbourhood_labels,
    aes(x = longitude, y = latitude,
        label = borough), label.size = 0.5) +
  labs(x = "", y = "") +
  theme_bw() +
  theme(axis.text.x = element_blank(),
        axis.text.y = element_blank(),
        axis.ticks = element_blank())

save(map_plot, file = "Report Figures/map_plot.rds")

```

Map Plot